

script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using Sagemaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances


▼ Content preferences

Language

English ▼





[SageMaker Immersion Day](#) > [Lab 6. SageMaker Pipelines](#) > **Option 2: Batch Inference Pipeline (SageMaker Pipelines)**

Option 2: Batch Inference Pipeline (SageMaker Pipelines)

 Make sure you have performed the steps described in the **Prerequisites** section before beginning this lab.

- [Overview](#)
- [Prerequisites](#)
- [Register model into SageMaker Model Registry](#)
 1. [Upload Model Artifact to S3 Bucket](#)
 2. [Create Model Group](#)
 3. [Register Model in Model Registry](#)
 4. [Approve Model in Model Registry](#)
- [Build the pipeline components](#)
 1. [Import statements and declare parameters and constants](#)
 2. [Generate Data for Inferences](#)
 3. [Upload Inferences Data to S3 Bucket](#)
 4. [Info about the Trained Model \(An Approved ModelPackage in SageMaker Model Registry\)](#)
 5. [Define create model step](#)
 6. [Define Transform Step to Perform Batch Transformation](#)
- [Build and Trigger the pipeline run](#)
- [Conclusion](#)

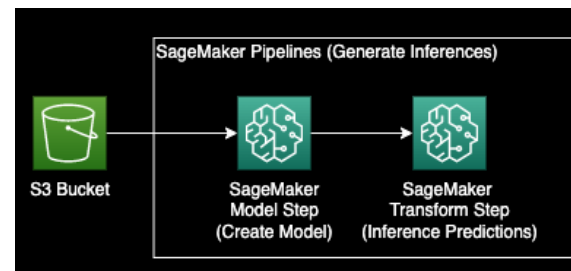
Overview

[Amazon SageMaker Pipelines](#) , is a capability of [Amazon SageMaker](#)  that makes it easy for data scientists and engineers to build, automate, and scale end to end machine learning pipelines. SageMaker Pipelines is a native [workflow orchestration tool](#)  for building ML pipelines that take advantage of direct [Amazon SageMaker](#)  integration.

This lab focuses on integration of offline hosting or batch transformation capability of with SageMaker Pipelines to generate predictions using a customer churn classification example. This pipeline includes steps to create model in SageMaker and generate inferences using SageMaker Batch Transform.

- script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab 3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language

The following diagram illustrates the high-level architecture of the ML workflow with the different steps to generate inferences using the trained model artifacts.



Inference Pipeline consists of the following steps:

1. Create a model in SageMaker using the latest approved model from SageMaker Model Registry.
2. Generate Inferences using the trained model artifacts.

Prerequisites

- Follow the instructions to [launch Amazon SageMaker Studio](#)
- Please ensure that you have git cloned the [repository](#) in your SageMaker Studio environment.

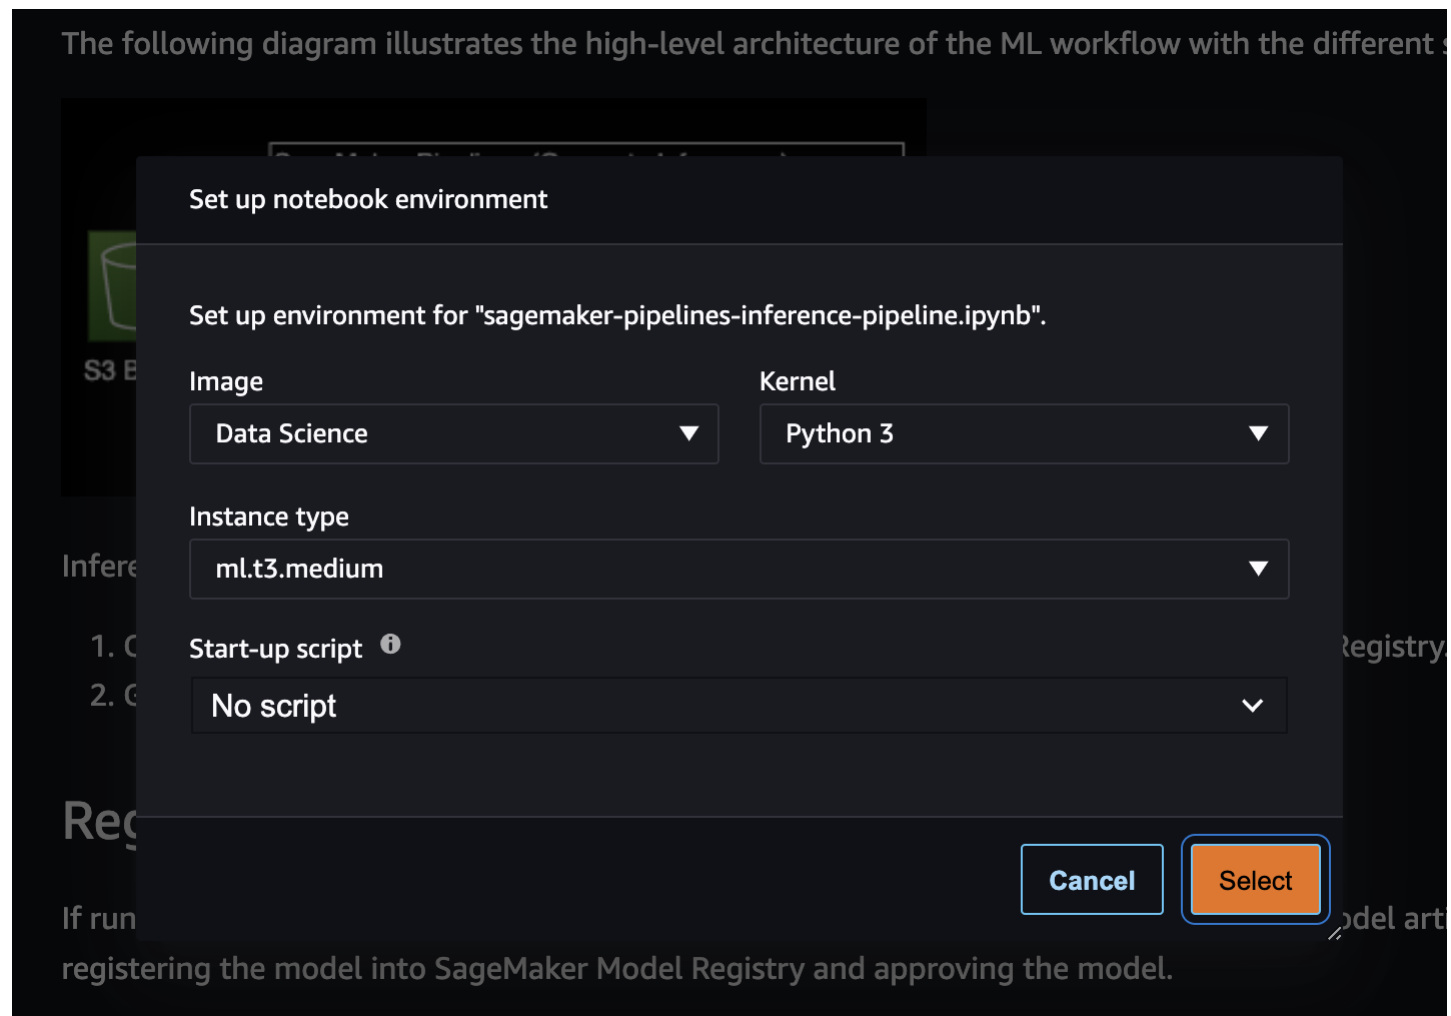
Register model into SageMaker Model Registry

If running through this lab independently, go through the optional step of uploading the model artifact [customer-retention-model.tar.gz](#) into S3 Bucket, registering the model into SageMaker Model Registry and approving the model.

Click on the “amazon-sagemaker-immersion-day” folder and then double click on the **sagemaker-pipelines-inference-pipeline.ipynb** notebook within the project home directory.

If you are prompted to choose a Kernel, choose the “Python 3 (Data Science)” kernel and click “Select”.

- script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language



Upload Model Artifact to S3 Bucket

Run the first cell within the notebook to set up SageMaker and S3 client objects and set up the S3 bucket location using the default bucket that comes with a SageMaker session:

Script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab 3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data



Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using SageMaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances

▼ Content preferences

Language

Upload Model Artifact to S3 Bucket

```
[2]: import boto3
import sagemaker

sagemaker_session = sagemaker.session.Session()
default_bucket = sagemaker_session.default_bucket()
s3_client = boto3.resource('s3')
s3_client.Bucket(default_bucket).upload_file("model/customer-retention-model.tar.gz", "churn/model_artifacts/customer-retention-model.tar.gz")
```

Create Model Group

Run the next cell which within SageMaker Model Registry, creates a [model group](#) that will track a group of versioned models.

Create Model Group

```
[3]: import time
import os
```



akadam ▼

```
region = sagemaker_session.default_region
sm_client = boto3.client('sagemaker', region_name=region)
model_package_group_input_dict = {
    "ModelPackageGroupName": model_package_group_name,
}

create_model_package_group_response = sm_client.create_model_package_group(**model_package_group_input_dict)
print('ModelPackageGroup Arn : {}'.format(create_model_package_group_response['ModelPackageGroupArn']))
```

Register Model in Model Registry

In the next cell, you can register a model into SageMaker Model Registry, by specifying the containers, model artifact and the associated environment variables.

Script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab 3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Register Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using SageMaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances

▼ Content preferences

Language

Register Model in Model Registry

```
[4]: # retrieve the image uri used to train model
image_uri = sagemaker.image_uris.retrieve(
    framework="xgboost",
    region=region,
    version="1.0-1",
    py_version="py3"
)

# Specify the model source
model_url = f"s3://{default_bucket}/churn/model_artifacts/customer-retention-model.tar.gz"

modelpackage_inference_specification = {
    "InferenceSpecification": {
        "Containers": [
            {
                "Image": image_uri,
                "ModelDataUrl": model_url
```

[Privacy policy](#) [Terms of use](#)

```
        "SupportedContentTypes": [ "text/csv" ],
        "SupportedResponseMIMETypes": [ "text/csv" ],
    }
}

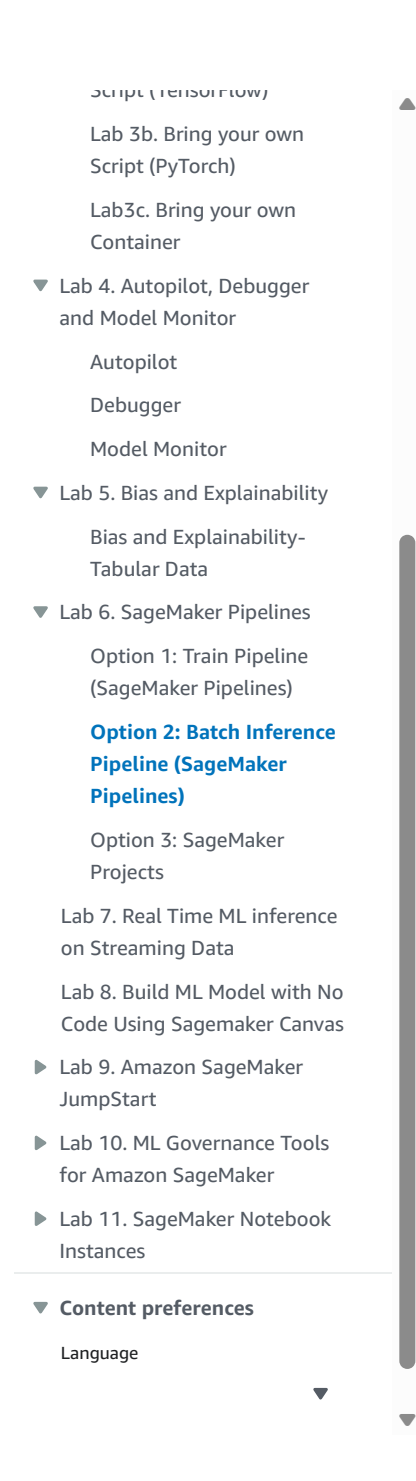
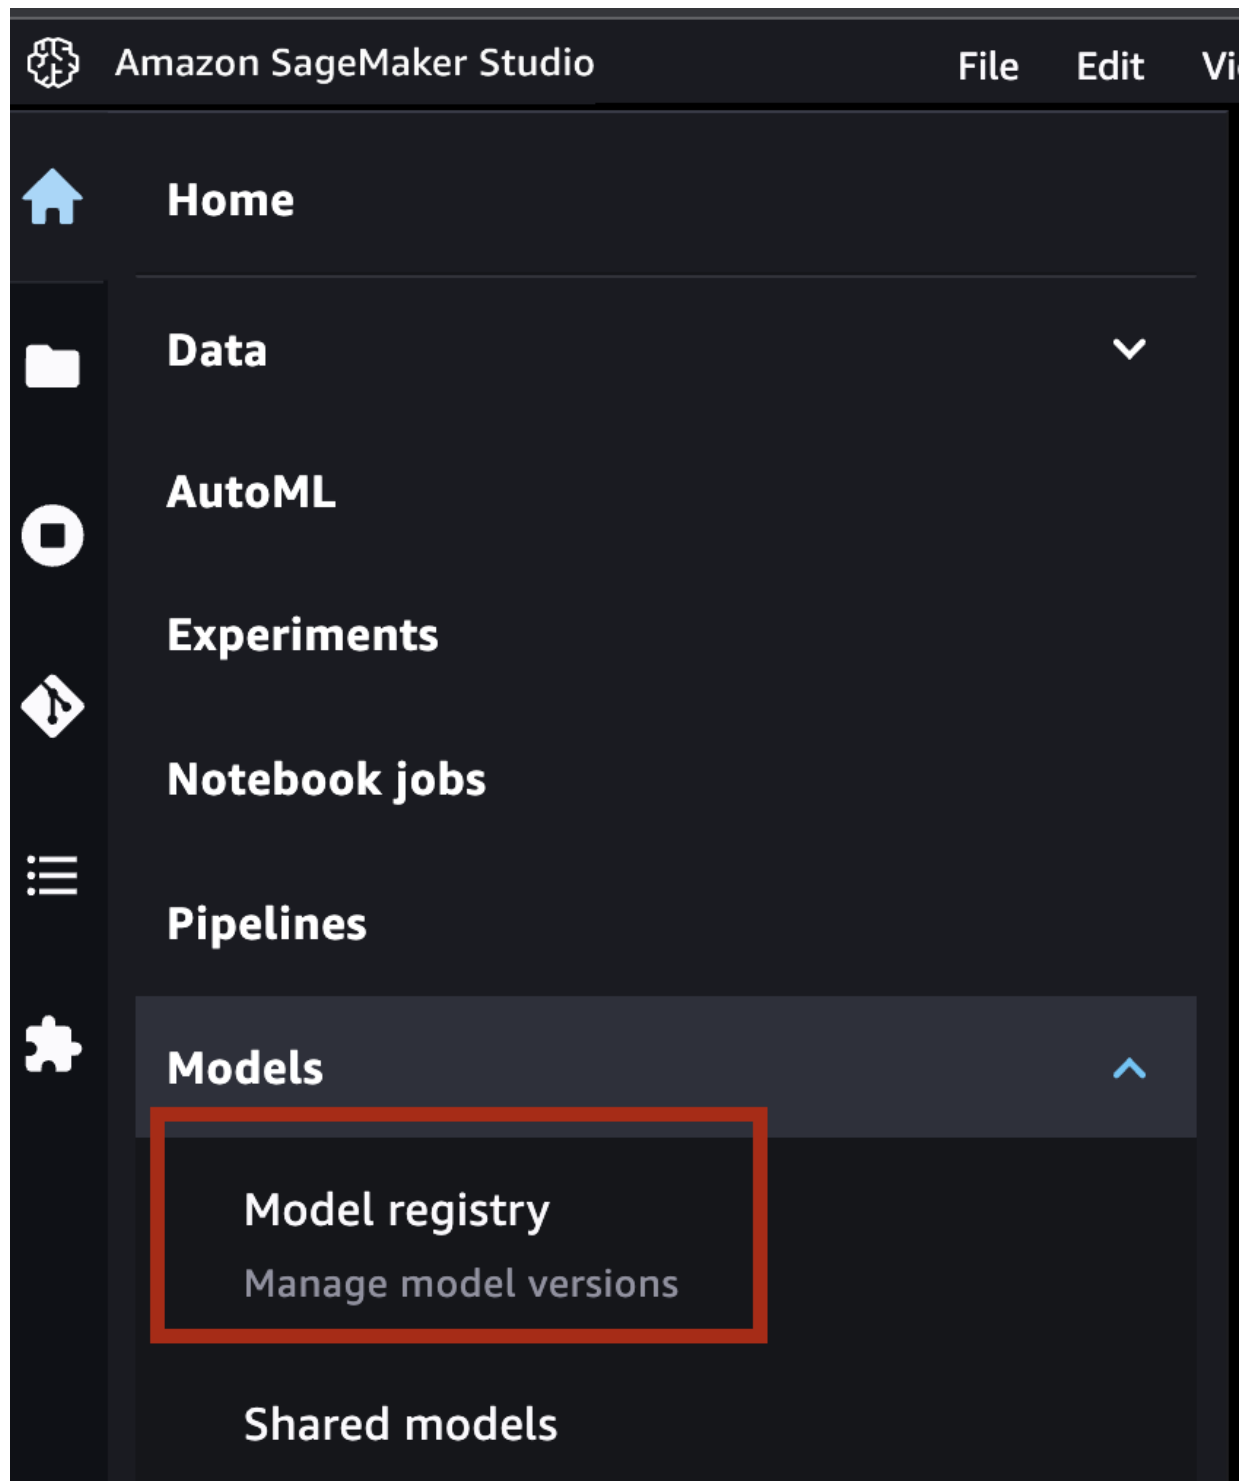
# Alternatively, you can specify the model source like this:
# modelpackage_inference_specification["InferenceSpecification"]["Containers"][0]["ModelDataUrl"]=model_url

create_model_package_input_dict = {
    "ModelPackageGroupName" : model_package_group_name,
    "ModelPackageDescription" : "Model to detect 3 different types of irises (Setosa, Versicolour, and Virginica)",
    "ModelApprovalStatus" : "PendingManualApproval"
}
create_model_package_input_dict.update(modelpackage_inference_specification)

create_model_package_response = sm_client.create_model_package(**create_model_package_input_dict)
model_package_arn = create_model_package_response["ModelPackageArn"]
print('ModelPackage Version ARN : {}'.format(model_package_arn))
```

Approve Model in Model Registry

The model registered within model registry can be checked by going to the home screen and choosing the *Models* → *Model Registry*.



Manage shared models & notebooks

script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab 3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using SageMaker Canvas

► Lab 9. Amazon SageMaker JumpStart

► Lab 10. ML Governance Tools for Amazon SageMaker

► Lab 11. SageMaker Notebook Instances

▼ Content preferences

Language

you can click on the Update Status tab and manually approve the model.

The screenshot shows the Amazon SageMaker Studio interface. The left sidebar contains a navigation menu with options like Home, Data, AutoML, Experiments, Notebook jobs, Pipelines, and Models. The 'Models' section is expanded, showing 'Model registry', 'Shared models', 'Inference compiler', and 'Edge packager'. The main panel displays the 'Model registry' for 'VERSION 1'. It includes a table with columns for Status, Pipeline, Execution, and Model group. The 'Status' column shows 'Pending'. The 'Update status' button is highlighted with a red box. Below this, there are tabs for Activity, Model quality, Explainability, Bias report, Inference recommender, Load test, and Settings. The 'Settings' tab is active, showing a table with 'Info' and 'Value' columns. The 'Metadata' section shows details like Project, Pipeline, Execution, Model group, Modified on, Modified by, Created on, Created by, and ARN.

Info	Value
Name	—
Description	—
Tags	—

Metadata	Value
Project	No project
Pipeline	sagemaker-immersion-train-pipeline
Execution	execution-1675035730036
Model group	churn-job-model-packages
Modified on	—
Modified by	—
Created on	49 minutes ago
Created by	—
ARN	arn:aws:sagemaker:us-west-2:123456789012:model-package/churn-job-model-packages/1
Transform instance types	ml.m5.xlarge
Realtime inference instance types	ml.t2.medium, ml.m5.xlarge
Trial Component	47gmogm6ctwe-ChurnHy-qWwhLCGzw4-001-0e378bb6-aws-training-job

Build the pipeline components

Import statements and declare parameters and constants

Run the next cell to set up SageMaker and S3 client objects, create PipelineSession, and set up the S3 bucket location using the default bucket that comes with a SageMaker session:

- script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language

Step 1: Import statements and declare parameters and constants

```
•[5]: import boto3
import pandas as pd
import sagemaker
from sagemaker.workflow.pipeline_context import PipelineSession

s3_client = boto3.resource('s3')
pipeline_name = f"sagemaker-immersion-inference-pipeline"
sagemaker_session = sagemaker.session.Session()
region = sagemaker_session.boto_region_name
role = sagemaker.get_execution_role()
pipeline_session = PipelineSession()
default_bucket = sagemaker_session.default_bucket()
model_package_group_name = f"ChurnModelPackageGroup"
```

Pipelines supports parameterization, which allows you to specify input parameters at runtime without changing your pipeline code. You can use the modules available under the `sagemaker.workflow.parameters` module, such as `ParameterInteger`, `ParameterFloat`, and `ParameterString`, to specify pipeline parameters of various data types. Run the following code to set up multiple input parameters:

```
[6]: from sagemaker.workflow.parameters import (
    ParameterInteger,
    ParameterString,
    ParameterFloat)

base_job_prefix = "churn-example"
processing_instance_count = ParameterInteger(name="ProcessingInstanceCount", default_value=1)
processing_instance_type = ParameterString(name="ProcessingInstanceType", default_value="ml.m5.xlarge")
transform_instance_type = ParameterString(name="TransformInstanceType", default_value="ml.m5.xlarge")
transform_instance_count = ParameterInteger(name="TransformInstanceCount", default_value=1)
batch_data_path = "s3://{}/data/batch/batch.csv".format(default_bucket)
model_approval_status = ParameterString(name="ModelApprovalStatus", default_value="PendingManualApproval")
```

Generate Data for Inferences

Run the next cell to generate the batch dataset, which you use later in the batch transform step. For this, download and save the [sample dataset](#) if not already done into the project folder. Sample on the original dataset to generate the data although in the real world scenario this would be a new dataset which the model has not been trained on before.

- Script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language

```

•[9]: def preprocess_batch_data(file_path):
    df = pd.read_csv(file_path)
    ## Convert to datetime columns
    df["firstorder"] = pd.to_datetime(df["firstorder"], errors='coerce')
    df["lastorder"] = pd.to_datetime(df["lastorder"], errors='coerce')
    ## Drop Rows with null values
    df = df.dropna()
    ## Create Column which gives the days between the last order and the first order
    df["first_last_days_diff"] = (df["lastorder"] - df["firstorder"]).dt.days
    ## Create Column which gives the days between when the customer record was created and the first order
    df["created"] = pd.to_datetime(df["created"])
    df["created_first_days_diff"] = (df["created"] - df["firstorder"]).dt.days
    ## Drop Columns
    df.drop(['custid', 'created', 'firstorder', 'lastorder'], axis=1, inplace=True)
    ## Apply one hot encoding on favday and city columns
    df = pd.get_dummies(df, prefix=['favday', 'city'], columns=['favday', 'city'])
    return df

# convert the store_data file into csv format
store_data = pd.read_excel("storedata_total.xlsx")
store_data.to_csv("storedata_total.csv")

[10]: # preprocess batch data and save into the data folder
batch_data = preprocess_batch_data("storedata_total.csv")
batch_data.pop("retained")
batch_sample = batch_data.sample(frac=0.2)
pd.DataFrame(batch_sample).to_csv("batch.csv", header=False, index=False)

```

Upload Inferences Data to S3 Bucket

Upload the dataset to Amazon S3.

Step 3: Upload Inferences Data to S3 Bucket

```

[11]: s3_client.Bucket(default_bucket).upload_file("batch.csv", "data/batch/batch.csv")

```

Info about the Trained Model

Retrieve the information of the model approved in SageMaker Model Registry using which the inferences will be generated.

- Script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language

```
[13]: sm_client = boto3.client("sagemaker")

# get a list of approved model packages from the model package group you specified earlier
approved_model_packages = sm_client.list_model_packages(
    ModelApprovalStatus='Approved',
    ModelPackageGroupName=model_package_group_name,
    SortBy='CreationTime',
    SortOrder='Descending'
)

# find the latest approved model package
try:
    latest_approved_model_package_arn = approved_model_packages['ModelPackageSummaryList'][0]['ModelPackageArn']
except Exception as e:
    print("Failed to retrieve an approved model package:", e)

print(latest_approved_model_package_arn)

# retrieve required information about the model
latest_approved_model_package_descr = sm_client.describe_model_package(ModelPackageName = latest_approved_model_package_arn)

# model artifact uri (tar.gz file)
model_artifact_uri = latest_approved_model_package_descr['InferenceSpecification']['Containers'][0]['ModelDataUrl']
# sagemaker image in ecr
image_uri = latest_approved_model_package_descr['InferenceSpecification']['Containers'][0]['Image']
```

Define create model step

Run the next cell to create a SageMaker model using the Pipelines model step. This step utilizes the approved model from SageMaker Model Registry.

```
[15]: from sagemaker import Model
      from sagemaker.inputs import CreateModelInput
      from sagemaker.workflow.model_step import ModelStep

      model = Model(
          image_uri=image_uri,
          model_data=model_artifact_uri,
          sagemaker_session=pipeline_session,
          role=role
      )

      step_create_model = ModelStep(
          name="ChurnCreateModel",
          step_args=model.create(instance_type="ml.m5.large", accelerator_type="ml.eia1.medium"),
      )
```

Define Transform Step to Perform Batch Transformation

- script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using SageMaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ Content preferences
 - Language

Run the next cell to run batch transformation using the trained model with the batch input created in the first step:

```
[16]: from sagemaker.transformer import Transformer
      from sagemaker.inputs import TransformInput
      from sagemaker.workflow.steps import TransformStep

      transformer = Transformer(
          model_name=step_create_model.properties.ModelName,
          instance_type="ml.m5.xlarge",
          instance_count=1,
          output_path=f"s3://{default_bucket}/ChurnTransform",
          sagemaker_session=pipeline_session
      )

      step_transform = TransformStep(
          name="ChurnTransform",
          step_args=transformer.transform(
              data=batch_data_path,
              content_type="text/csv"
          )
      )
```

Build and Trigger the pipeline run

After defining all of the component steps, you can assemble them into a Pipelines object. You don't need to specify the order of pipeline because Pipelines automatically infers the order sequence based on the dependencies between the steps.

script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using SageMaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances

▼ Content preferences

Language

```
[17]: from sagemaker.workflow.pipeline import Pipeline

pipeline = Pipeline(
    name=pipeline_name,
    parameters=[
        processing_instance_type,
        processing_instance_count,
        transform_instance_type,
        transform_instance_count,
        batch_data,
    ],
    steps=[step_create_model, step_transform],
)
```

Run the following code in the next cell in your notebook. If the pipeline already exists, the code updates the pipeline. If the pipeline doesn't exist, it creates a new one.

```
[18]: # Create a new or update existing Pipeline
pipeline.upsert(role_arn=role)
# start Pipeline execution
pipeline.start()
```

Go to home within the studio domain and click on pipelines.



Home



Data



AutoML



Experiments



Notebook jobs



Pipelines



Models



Deployments



Quick start solutions



script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using Sagemaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances

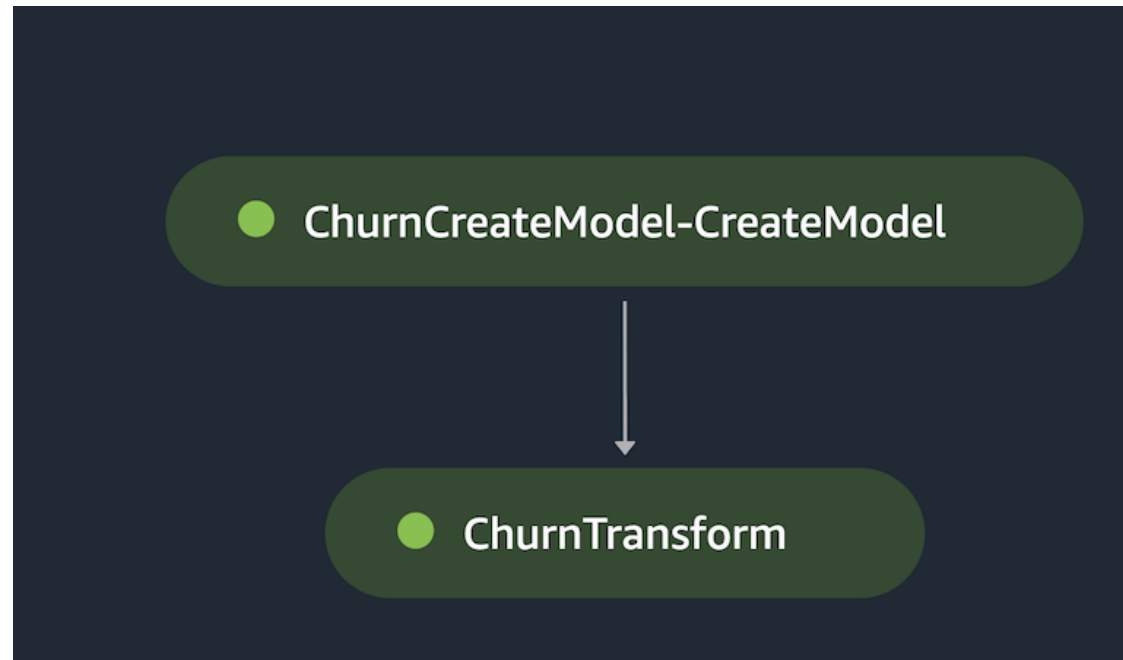
▼ Content preferences

Language

Learning resources



Once the pipeline execution is complete, the inferences generated should be available within the S3 Bucket.



Amazon S3 > Buckets > sagemaker-us-west-2- > ChurnTransform/

ChurnTransform/

Objects Properties

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

[Refresh](#) [Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions](#) [Create folder](#) [Upload](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	batch.csv.out	out	22:44:22 (UTC-07:00)	116.1 KB	Standard

script (TensorFlow)

Lab 3b. Bring your own
Script (PyTorch)

Lab3c. Bring your own
Container

▼ Lab 4. Autopilot, Debugger
and Model Monitor

Autopilot

Debugger

Model Monitor

▼ Lab 5. Bias and Explainability

Bias and Explainability-
Tabular Data

▼ Lab 6. SageMaker Pipelines

Option 1: Train Pipeline
(SageMaker Pipelines)

**Option 2: Batch Inference
Pipeline (SageMaker
Pipelines)**

Option 3: SageMaker
Projects

Lab 7. Real Time ML inference
on Streaming Data

Lab 8. Build ML Model with No
Code Using SageMaker Canvas

► Lab 9. Amazon SageMaker
JumpStart

► Lab 10. ML Governance Tools
for Amazon SageMaker

► Lab 11. SageMaker Notebook
Instances

▼ Content preferences

Language

- script (TensorFlow)
- Lab 3b. Bring your own Script (PyTorch)
- Lab3c. Bring your own Container
- ▼ Lab 4. Autopilot, Debugger and Model Monitor
 - Autopilot
 - Debugger
 - Model Monitor
- ▼ Lab 5. Bias and Explainability
 - Bias and Explainability-Tabular Data
- ▼ Lab 6. SageMaker Pipelines
 - Option 1: Train Pipeline (SageMaker Pipelines)
 - Option 2: Batch Inference Pipeline (SageMaker Pipelines)**
 - Option 3: SageMaker Projects
- Lab 7. Real Time ML inference on Streaming Data
- Lab 8. Build ML Model with No Code Using Sagemaker Canvas
- Lab 9. Amazon SageMaker JumpStart
- Lab 10. ML Governance Tools for Amazon SageMaker
- Lab 11. SageMaker Notebook Instances
- ▼ **Content preferences**
 - Language

Conclusion

In this lab we have walked through how to build an ML workflow with the different steps to generate inferences using [batch transform](#) with trained model artifact along with other built-in SageMaker features for churn prediction. The solution can be extended with other steps as needed to implement your own ML workflow. To learn more about SageMaker Pipelines, check out the [website](#) and the [documentation](#).

[Previous](#)[Next](#)